

# CSC2/458 Parallel and Distributed Systems Machines and Models

---

Sreepathi Pai

January 23, 2018

URCS

# Outline

Recap

Scalability

Taxonomy of Parallel Machines

Performance Metrics

# Outline

Recap

Scalability

Taxonomy of Parallel Machines

Performance Metrics

# Goals

What is the goal of parallel programming?

# Scalability

Why is scalability important?

# Outline

Recap

Scalability

Taxonomy of Parallel Machines

Performance Metrics

# Speedup

$$\text{Speedup}(n) = \frac{T_1}{T_n}$$

- $T_1$  is time on one processor
- $T_n$  is time on  $n$  processors

# Amdahl's Law

Let:

- $T_1$  be  $T_{serial} + T_{parallelizable}$
- $T_n$  is then  $T_{serial} + \frac{T_{parallelizable}}{n}$ , assuming perfect scalability

Divide both terms  $T_1$  and  $T_n$  by  $T_1$  to obtain serial and parallelizable ratios.

$$\text{Speedup}(n) = \frac{1}{r_{serial} + \frac{r_{parallelizable}}{n}}$$



## Amdahl's Law – In the limit

$$\text{Speedup}(\infty) = \frac{1}{r_{\text{serial}}}$$

This is also known as *strong scalability* – work is fixed and number of processors is varied.

What are the implications of this?

# Scalability Limits

Assuming infinite processors, what is the speedup if:

- serial ratio  $r_{serial}$  is 0.5 (i.e. 50%)
- serial ratio is 0.1 (i.e. 10%)
- serial ratio is 0.01 (i.e. 1%)

## Current Top 5 supercomputers

- Sunway TaihuLight (10.6M cores)
- Tianhe 2 (3.1M cores)
- Piz Daint (361K cores)
- Gyoukou (19.8M cores)
- Titan (560K cores)

Source: Top 500

# Weak Scalability

- Work increases as number of processors increase
  - Parallel work should increase linearly with processors
- Work  $W = \alpha W + (1 - \alpha)W$ 
  - $\alpha$  is serial fraction of work
- Scaled Work  $W' = \alpha W + n(1 - \alpha)W$ 
  - Empirical observation
  - Usually referred to as Gustafson's Law

Source:

<http://www.johngustafson.net/pubs/pub13/amdahl.htm>

# Outline

Recap

Scalability

Taxonomy of Parallel Machines

Performance Metrics

# Organization of Parallel Computers

Components of parallel machines:

- Processing Elements
- Memories
- Interconnect
  - how processors, memories are connected to each other

# Flynn's Taxonomy

- Based on notion of “streams”
  - Instruction stream
  - Data stream
- Taxonomy based on number of each type of streams
  - Single Instruction - Single Data (SISD)
  - Single Instruction - Multiple Data (SIMD)
  - Multiple Instruction - Single Data (MISD)
  - Multiple Instruction - Multiple Data (MIMD)


Flynn, J., (1966), “<http://ieeexplore.ieee.org/document/1447203>/Very High Speed Computing Systems”, Proceedings of the IEEE

# SIMD Implementations: Vector Machines

The Cray-1 (circa 1977):


- $V_x$  – vector registers
  - 64 elements
  - 64-bits per element
- Vector length register ( $Vlen$ )
- Vector mask register

Richard Russell, "The Cray-1 Computer System", Comm. ACM 21,1 (Jan 1978), 63-72





## Vector Instructions – Vertical



For  $0 < i < Vlen$ :

$$dst[i] = src1[i] + src2[i]$$

- Most arithmetic instructions

## Vector Instructions – Horizontal

$$1 = \min( \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline \end{array} )$$


For  $0 < i < Vlen$ :

`dst = min(src1[i], dst)`

Note that `dst` is a scalar.

- Mostly reductions (min, max, sum, etc.)
- Not well supported
  - Cray-1 did not have this

## Vector Instructions – Shuffle/Permute



```
dst = shuffle(src1, mask)
```

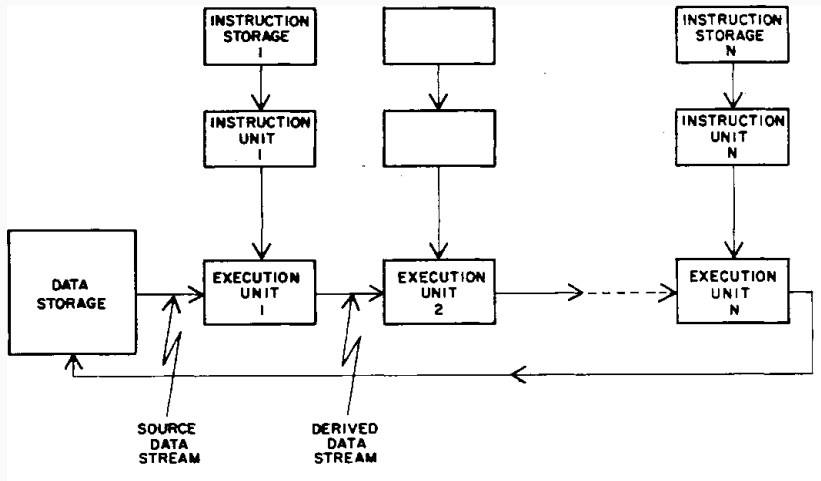
- Poor support on older implementations
- Reasonably well-supported on recent implementations

# Masking/Predication

src1	6	5	7	2
g5mask	1	0	1	0
src1	6	5	7	2
*				
src2	1	4	2	2
=				
dst	6	?	14	?


```
g5mask = gt(src1, 5)
dst = mul(src1, src2, g5mask)
```

# MISD - ?



Flynn, J., (1966), "<http://ieeexplore.ieee.org/document/1447203>/Very High Speed Computing Systems", Proceedings of the IEEE

# What type of machine is this? Hyperthreaded Core



Different colours in RAM indicate different instruction streams.

Source: <https://en.wikipedia.org/wiki/Hyper-threading>


# What type of machine is this? GPU



Each instruction is 32-wide.

Source: <https://devblogs.nvidia.com/inside-pascal/>

# What type of machine is this? TPU Matrix Multiply Unit

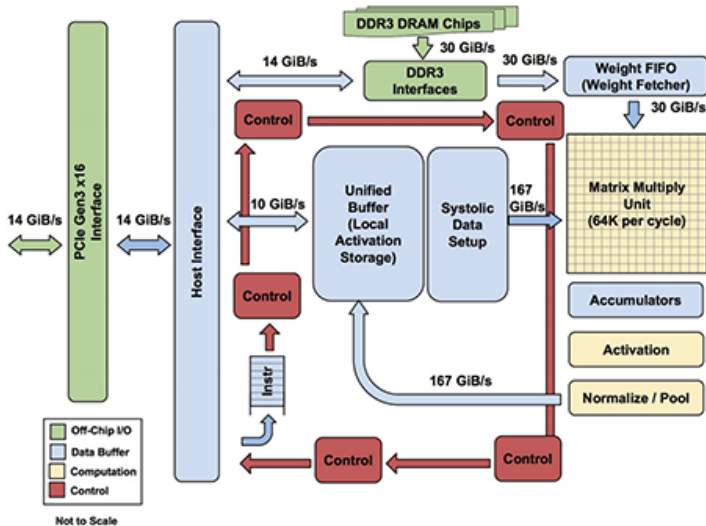


Source:

<https://cloud.google.com/blog/big-data/2017/05/an-in-depth-look-at-googles-first-tensor-processing-unit-tpu>



# TPU Overview



Source:

<https://cloud.google.com/blog/big-data/2017/05/an-in-depth-look-at-googles-first-tensor-processing-unit-tpu>

# Modern Multicores

- Multiple Cores (MIMD)
- (Short) Vector Instruction Sets (SIMD)
  - MMX, SSE, AVX (Intel)
  - 3DNow (AMD)
  - NEON (ARM)

# Outline

Recap

Scalability

Taxonomy of Parallel Machines

Performance Metrics

# Metrics we care about

- Latency
  - Time to complete task
  - Lower is better
- Throughput
  - Rate of completing tasks
  - Higher is better
- Utilization
  - Time “worker” (processor, unit) is busy
  - Higher is better
- Speedup
  - Higher is better

# Reducing Latency

- Use cheap operations
- Which of these operations are expensive?
  - Bitshift
  - Integer Divide
  - Integer Multiply
- Latency fundamentally bounded by physics

# Increasing Throughput

- Parallelize!
  - Lots of techniques, focus of this class
- Add more processors
- Need lots of work though to benefit

# Speedup

- Measure speedup w.r.t. fastest serial code
  - Not parallel program on 1 processor
- Always report runtime
  - Never speedup alone
- Are superlinear speedups possible?